

dplyr tutorial

What is dplyr?

dplyr is a powerful R-package to transform and summarize tabular data with rows and columns. For another explanation of dplyr see the dplyr package vignette: [Introduction to dplyr](#)

Why Is It Useful?

The package contains a set of functions (or “verbs”) that perform common data manipulation operations such as filtering for rows, selecting specific columns, re-ordering rows, adding new columns and summarizing data.

In addition, dplyr contains a useful function to perform another common task which is the “split-apply-combine” concept. We will discuss that in a little bit.

How Does It Compare To Using Base Functions R?

If you are familiar with R, you are probably familiar with base R functions such as `split()`, `subset()`, `apply()`, `sapply()`, `lapply()`, `tapply()` and `aggregate()`. Compared to base functions in R, the functions in dplyr are easier to work with, are more consistent in the syntax and are targeted for data analysis around data frames, instead of just vectors.

How Do I Get dplyr?

To install dplyr:

```
install.packages("dplyr")
```

To load dplyr:

```
library(dplyr)
```

Data: Mammals Sleep

The msleep (mammals sleep) data set contains the sleep times and weights for a set of mammals and is available in the dagdata repository on github. This data set contains 83 rows and 11 variables.

Download the msleep data set in CSV format from [here](#), and then load into R:

```
dir <- ("/Users/guillermotorres/Documents/Proyectos/2017/12PastoWorkshop/CourseLab/ws/extdata/") # defini
file <- paste0(dir,"msleep_ggplot2.csv")
msleep <- read.csv(file)
head(msleep)
```

##	name	genus	vore	order	conservation
## 1	Cheetah	Acinonyx	carni	Carnivora	1c
## 2	Owl monkey	Aotus	omni	Primates	<NA>
## 3	Mountain beaver	Aplodontia	herbi	Rodentia	nt
## 4	Greater short-tailed shrew	Blarina	omni	Soricomorpha	1c
## 5	Cow	Bos	herbi	Artiodactyla	domesticated
## 6	Three-toed sloth	Bradypus	herbi	Pilosa	<NA>

```
##   sleep_total sleep_rem sleep_cycle awake brainwt  bodywt
## 1         12.1        NA         NA  11.9        NA  50.000
## 2         17.0         1.8         NA   7.0 0.01550   0.480
## 3         14.4         2.4         NA   9.6        NA   1.350
## 4         14.9         2.3 0.1333333   9.1 0.00029   0.019
## 5          4.0         0.7 0.6666667  20.0 0.42300 600.000
## 6         14.4         2.2 0.7666667   9.6        NA   3.850
```

The columns (in order) correspond to the following:

column name	Description
name	common name
genus	taxonomic rank
vore	carnivore, omnivore or herbivore?
order	taxonomic rank
conservation	the conservation status of the mammal
sleep_total	total amount of sleep, in hours
sleep_rem	rem sleep, in hours
sleep_cycle	length of sleep cycle, in hours
awake	amount of time spent awake, in hours
brainwt	brain weight in kilograms
bodywt	body weight in kilograms

Important dplyr Verbs To Remember

dplyr verbs	Description
<code>select()</code>	select columns
<code>filter()</code>	filter rows
<code>arrange()</code>	re-order or arrange rows
<code>mutate()</code>	create new columns
<code>summarise()</code>	summarise values
<code>group_by()</code>	allows for group operations in the “split-apply-combine” concept

dplyr Verbs In Action

The two most basic functions are `select()` and `filter()`, which selects columns and filters rows respectively.

Selecting Columns Using `select()`

Select a set of columns: the name and the sleep_total columns.

```
sleepData <- select(msleep, name, sleep_total)
head(sleepData)
```

```
##           name sleep_total
## 1      Cheetah         12.1
## 2 Owl monkey         17.0
## 3 Mountain beaver        14.4
## 4 Greater short-tailed shrew 14.9
```

```
## 5          Cow          4.0
## 6    Three-toed sloth    14.4
```

To select all the columns *except* a specific column, use the “-” (subtraction) operator (also known as negative indexing):

```
head(select(msleep, -name))
```

```
##      genus  vore      order conservation sleep_total sleep_rem
## 1  Acinonyx carni   Carnivora          lc         12.1      NA
## 2    Aotus  omni    Primates        <NA>         17.0      1.8
## 3 Aplodontia herbi   Rodentia          nt         14.4      2.4
## 4   Blarina  omni Soricomorpha          lc         14.9      2.3
## 5      Bos  herbi Artiodactyla domesticated         4.0      0.7
## 6  Bradypus herbi    Pilosa        <NA>         14.4      2.2
##  sleep_cycle awake brainwt  bodywt
## 1          NA  11.9      NA  50.000
## 2          NA   7.0 0.01550   0.480
## 3          NA   9.6      NA   1.350
## 4  0.1333333   9.1 0.00029   0.019
## 5  0.6666667  20.0 0.42300 600.000
## 6  0.7666667   9.6      NA   3.850
```

To select a range of columns by name, use the “:” (colon) operator:

```
head(select(msleep, name:order))
```

```
##      name      genus  vore      order
## 1   Cheetah  Acinonyx carni   Carnivora
## 2 Owl monkey    Aotus  omni    Primates
## 3 Mountain beaver Aplodontia herbi   Rodentia
## 4 Greater short-tailed shrew   Blarina  omni Soricomorpha
## 5      Cow      Bos  herbi Artiodactyla
## 6 Three-toed sloth  Bradypus herbi    Pilosa
```

To select all columns that start with the character string “sl”, use the function `starts_with()`:

```
head(select(msleep, starts_with("sl")))
```

```
##  sleep_total sleep_rem sleep_cycle
## 1         12.1      NA          NA
## 2         17.0      1.8          NA
## 3         14.4      2.4          NA
## 4         14.9      2.3  0.1333333
## 5          4.0      0.7  0.6666667
## 6         14.4      2.2  0.7666667
```

Some additional options to select columns based on a specific criteria include:

1. `ends_with()` = Select columns that end with a character string
2. `contains()` = Select columns that contain a character string
3. `matches()` = Select columns that match a regular expression
4. `one_of()` = Select column names that are from a group of names

Selecting Rows Using `filter()`

Filter the rows for mammals that sleep a total of more than 16 hours.

```
filter(msleep, sleep_total >= 16)
```

```
##           name      genus  vore      order conservation
## 1      Owl monkey      Aotus  omni      Primates      <NA>
## 2 Long-nosed armadillo  Dasypus  carni      Cingulata      lc
## 3 North American Opossum  Didelphis  omni  Didelphimorphia      lc
## 4      Big brown bat  Eptesicus  insecti      Chiroptera      lc
## 5 Thick-tailed opossum  Lutreolina  carni  Didelphimorphia      lc
## 6 Little brown bat      Myotis  insecti      Chiroptera      <NA>
## 7      Giant armadillo  Priodontes  insecti      Cingulata      en
## 8 Arctic ground squirrel  Spermophilus  herbi      Rodentia      lc
##  sleep_total sleep_rem sleep_cycle awake brainwt bodywt
## 1      17.0      1.8      NA      7.0 0.01550 0.480
## 2      17.4      3.1 0.3833333 6.6 0.01080 3.500
## 3      18.0      4.9 0.3333333 6.0 0.00630 1.700
## 4      19.7      3.9 0.1166667 4.3 0.00030 0.023
## 5      19.4      6.6      NA      4.6      NA 0.370
## 6      19.9      2.0 0.2000000 4.1 0.00025 0.010
## 7      18.1      6.1      NA      5.9 0.08100 60.000
## 8      16.6      NA      NA      7.4 0.00570 0.920
```

Filter the rows for mammals that sleep a total of more than 16 hours *and* have a body weight of greater than 1 kilogram.

```
filter(msleep, sleep_total >= 16, bodywt >= 1)
```

```
##           name      genus  vore      order conservation
## 1 Long-nosed armadillo  Dasypus  carni      Cingulata      lc
## 2 North American Opossum  Didelphis  omni  Didelphimorphia      lc
## 3      Giant armadillo  Priodontes  insecti      Cingulata      en
##  sleep_total sleep_rem sleep_cycle awake brainwt bodywt
## 1      17.4      3.1 0.3833333 6.6 0.0108 3.5
## 2      18.0      4.9 0.3333333 6.0 0.0063 1.7
## 3      18.1      6.1      NA      5.9 0.0810 60.0
```

Filter the rows for mammals in the Perissodactyla and Primates taxonomic order

```
filter(msleep, order %in% c("Perissodactyla", "Primates"))
```

```
##           name      genus  vore      order conservation
## 1      Owl monkey      Aotus  omni      Primates      <NA>
## 2      Grivet Cercopithecus  omni      Primates      lc
## 3      Horse      Equus  herbi  Perissodactyla  domesticated
## 4      Donkey      Equus  herbi  Perissodactyla  domesticated
## 5      Patas monkey  Erythrocebus  omni      Primates      lc
## 6      Galago      Galago  omni      Primates      <NA>
## 7      Human      Homo  omni      Primates      <NA>
## 8      Mongoose lemur      Lemur  herbi      Primates      vu
## 9      Macaque      Macaca  omni      Primates      <NA>
## 10 Slow loris      Nyctibeus  carni      Primates      <NA>
## 11 Chimpanzee      Pan  omni      Primates      <NA>
## 12 Baboon      Papio  omni      Primates      <NA>
## 13 Potto      Perodicticus  omni      Primates      lc
## 14 Squirrel monkey      Saimiri  omni      Primates      <NA>
## 15 Brazilian tapir      Tapirus  herbi  Perissodactyla      vu
##  sleep_total sleep_rem sleep_cycle awake brainwt bodywt
```

## 1	17.0	1.8	NA	7.0	0.0155	0.480
## 2	10.0	0.7	NA	14.0	NA	4.750
## 3	2.9	0.6	1.0000000	21.1	0.6550	521.000
## 4	3.1	0.4	NA	20.9	0.4190	187.000
## 5	10.9	1.1	NA	13.1	0.1150	10.000
## 6	9.8	1.1	0.5500000	14.2	0.0050	0.200
## 7	8.0	1.9	1.5000000	16.0	1.3200	62.000
## 8	9.5	0.9	NA	14.5	NA	1.670
## 9	10.1	1.2	0.7500000	13.9	0.1790	6.800
## 10	11.0	NA	NA	13.0	0.0125	1.400
## 11	9.7	1.4	1.4166667	14.3	0.4400	52.200
## 12	9.4	1.0	0.6666667	14.6	0.1800	25.235
## 13	11.0	NA	NA	13.0	NA	1.100
## 14	9.6	1.4	NA	14.4	0.0200	0.743
## 15	4.4	1.0	0.9000000	19.6	0.1690	207.501

You can use the boolean operators (e.g. `>`, `<`, `>=`, `<=`, `!=`, `%in%`) to create the logical tests.

Pipe Operator: `%>%`

Before we go any further, let's introduce the pipe operator: `%>%`. `dplyr` imports this operator from another package (`magrittr`). This operator allows you to pipe the output from one function to the input of another function. Instead of nesting functions (reading from the inside to the outside), the idea of piping is to read the functions from left to right.

Here's an example you have seen:

```
head(select(msleep, name, sleep_total))
```

##	name	sleep_total
## 1	Cheetah	12.1
## 2	Owl monkey	17.0
## 3	Mountain beaver	14.4
## 4	Greater short-tailed shrew	14.9
## 5	Cow	4.0
## 6	Three-toed sloth	14.4

Now in this case, we will pipe the `msleep` data frame to the function that will select two columns (`name` and `sleep_total`) and then pipe the new data frame to the function `head()`, which will return the head of the new data frame.

```
msleep %>%
  select(name, sleep_total) %>%
  head
```

##	name	sleep_total
## 1	Cheetah	12.1
## 2	Owl monkey	17.0
## 3	Mountain beaver	14.4
## 4	Greater short-tailed shrew	14.9
## 5	Cow	4.0
## 6	Three-toed sloth	14.4

You will soon see how useful the pipe operator is when we start to combine many functions.

Back To dplyr Verbs In Action

Now that you know about the pipe operator (`%>%`), we will use it throughout the rest of this tutorial.

Arrange Or Re-order Rows Using `arrange()`

To arrange (or re-order) rows by a particular column, such as the taxonomic order, list the name of the column you want to arrange the rows by:

```
msleep %>% arrange(order) %>% head
```

```
##      name      genus vore      order conservation sleep_total sleep_rem
## 1  Tenrec    Tenrec  omni Afrosoricida      <NA>         15.6         2.3
## 2    Cow      Bos   herbi Artiodactyla domesticated      4.0         0.7
## 3 Roe deer Capreolus herbi Artiodactyla      lc         3.0         NA
## 4    Goat    Capri herbi Artiodactyla      lc         5.3         0.6
## 5 Giraffe  Giraffa herbi Artiodactyla      cd         1.9         0.4
## 6  Sheep    Ovis   herbi Artiodactyla domesticated      3.8         0.6
##  sleep_cycle awake brainwt  bodywt
## 1          NA    8.4  0.0026   0.900
## 2  0.6666667  20.0  0.4230  600.000
## 3          NA   21.0  0.0982   14.800
## 4          NA   18.7  0.1150   33.500
## 5          NA   22.1      NA  899.995
## 6          NA   20.2  0.1750   55.500
```

Now we will select three columns from `msleep`, arrange the rows by the taxonomic order and then arrange the rows by `sleep_total`. Finally, show the head of the final data frame:

```
msleep %>%
  select(name, order, sleep_total) %>%
  arrange(order, sleep_total) %>%
  head
```

```
##      name      order sleep_total
## 1  Tenrec Afrosoricida      15.6
## 2 Giraffe Artiodactyla       1.9
## 3 Roe deer Artiodactyla       3.0
## 4  Sheep Artiodactyla       3.8
## 5    Cow Artiodactyla       4.0
## 6    Goat Artiodactyla       5.3
```

Same as above, except here we filter the rows for mammals that sleep for 16 or more hours, instead of showing the head of the final data frame:

```
msleep %>%
  select(name, order, sleep_total) %>%
  arrange(order, sleep_total) %>%
  filter(sleep_total >= 16)
```

```
##      name      order sleep_total
## 1 Big brown bat  Chiroptera      19.7
## 2 Little brown bat Chiroptera      19.9
## 3 Long-nosed armadillo Cingulata      17.4
## 4 Giant armadillo Cingulata      18.1
## 5 North American Opossum Didelphimorphia      18.0
```

```
## 6 Thick-tailed opossum Didelphimorphia 19.4
## 7 Owl monkey Primates 17.0
## 8 Arctic ground squirrel Rodentia 16.6
```

Something slightly more complicated: same as above, except arrange the rows in the sleep_total column in a descending order. For this, use the function desc()

```
msleep %>%
  select(name, order, sleep_total) %>%
  arrange(order, desc(sleep_total)) %>%
  filter(sleep_total >= 16)
```

```
##           name           order sleep_total
## 1 Little brown bat Chiroptera      19.9
## 2 Big brown bat Chiroptera      19.7
## 3 Giant armadillo Cingulata      18.1
## 4 Long-nosed armadillo Cingulata 17.4
## 5 Thick-tailed opossum Didelphimorphia 19.4
## 6 North American Opossum Didelphimorphia 18.0
## 7 Owl monkey Primates 17.0
## 8 Arctic ground squirrel Rodentia 16.6
```

Create New Columns Using mutate()

The mutate() function will add new columns to the data frame. Create a new column called rem_proportion, which is the ratio of rem sleep to total amount of sleep.

```
msleep %>%
  mutate(rem_proportion = sleep_rem / sleep_total) %>%
  head
```

```
##           name      genus vore      order conservation
## 1 Cheetah Acinonyx carni Carnivora lc
## 2 Owl monkey Aotus omni Primates <NA>
## 3 Mountain beaver Aplodontia herbi Rodentia nt
## 4 Greater short-tailed shrew Blarina omni Soricomorpha lc
## 5 Cow Bos herbi Artiodactyla domesticated
## 6 Three-toed sloth Bradypus herbi Pilosa <NA>
## sleep_total sleep_rem sleep_cycle awake brainwt bodywt rem_proportion
## 1 12.1 NA NA 11.9 NA 50.000 NA
## 2 17.0 1.8 NA 7.0 0.01550 0.480 0.1058824
## 3 14.4 2.4 NA 9.6 NA 1.350 0.1666667
## 4 14.9 2.3 0.1333333 9.1 0.00029 0.019 0.1543624
## 5 4.0 0.7 0.6666667 20.0 0.42300 600.000 0.1750000
## 6 14.4 2.2 0.7666667 9.6 NA 3.850 0.1527778
```

You can many new columns using mutate (separated by commas). Here we add a second column called bodywt_grams which is the bodywt column in grams.

```
msleep %>%
  mutate(rem_proportion = sleep_rem / sleep_total,
         bodywt_grams = bodywt * 1000) %>%
  head
```

```
##           name      genus vore      order conservation
## 1 Cheetah Acinonyx carni Carnivora lc
## 2 Owl monkey Aotus omni Primates <NA>
```

```
## 3      Mountain beaver Aplodontia herbi      Rodentia      nt
## 4 Greater short-tailed shrew Blarina omni Soricomorpha      lc
## 5              Cow      Bos herbi Artiodactyla domesticated
## 6      Three-toed sloth Bradypus herbi      Pilosa      <NA>
##  sleep_total sleep_rem sleep_cycle awake brainwt bodywt rem_proportion
## 1      12.1      NA      NA 11.9      NA 50.000      NA
## 2      17.0      1.8      NA 7.0 0.01550 0.480      0.1058824
## 3      14.4      2.4      NA 9.6      NA 1.350      0.1666667
## 4      14.9      2.3 0.1333333 9.1 0.00029 0.019      0.1543624
## 5      4.0      0.7 0.6666667 20.0 0.42300 600.000      0.1750000
## 6      14.4      2.2 0.7666667 9.6      NA 3.850      0.1527778
##  bodywt_grams
## 1      50000
## 2      480
## 3      1350
## 4      19
## 5      600000
## 6      3850
```

Create summaries of the data frame using summarise()

The `summarise()` function will create summary statistics for a given column in the data frame such as finding the mean. For example, to compute the average number of hours of sleep, apply the `mean()` function to the column `sleep_total` and call the summary value `avg_sleep`.

```
msleep %>%
  summarise(avg_sleep = mean(sleep_total))
```

```
##  avg_sleep
## 1  10.43373
```

There are many other summary statistics you could consider such `sd()`, `min()`, `max()`, `median()`, `sum()`, `n()` (returns the length of vector), `first()` (returns first value in vector), `last()` (returns last value in vector) and `n_distinct()` (number of distinct values in vector).

```
msleep %>%
  summarise(avg_sleep = mean(sleep_total),
            min_sleep = min(sleep_total),
            max_sleep = max(sleep_total),
            total = n())
```

```
##  avg_sleep min_sleep max_sleep total
## 1  10.43373      1.9      19.9     83
```

Group operations using group_by()

The `group_by()` verb is an important function in `dplyr`. As we mentioned before it's related to concept of "split-apply-combine". We literally want to split the data frame by some variable (e.g. taxonomic order), apply a function to the individual data frames and then combine the output.

Let's do that: split the `msleep` data frame by the taxonomic order, then ask for the same summary statistics as above. We expect a set of summary statistics for each taxonomic order.

```
msleep %>%
  group_by(order) %>%
```



```
summarise(avg_sleep = mean(sleep_total),
          min_sleep = min(sleep_total),
          max_sleep = max(sleep_total),
          total = n())
```

```
## # A tibble: 19 x 5
##       order avg_sleep min_sleep max_sleep total
##       <fctr>   <dbl>    <dbl>    <dbl> <int>
## 1 Afrosoricida 15.600000    15.6     15.6     1
## 2 Artiodactyla  4.516667     1.9      9.1     6
## 3 Carnivora    10.116667     3.5    15.8    12
## 4 Cetacea      4.500000     2.7     5.6     3
## 5 Chiroptera   19.800000    19.7    19.9     2
## 6 Cingulata    17.750000    17.4    18.1     2
## 7 Didelphimorphia 18.700000    18.0    19.4     2
## 8 Diprotodontia 12.400000    11.1    13.7     2
## 9 Erinaceomorpha 10.200000    10.1    10.3     2
## 10 Hyracoidea   5.666667     5.3     6.3     3
## 11 Lagomorpha   8.400000     8.4     8.4     1
## 12 Monotremata  8.600000     8.6     8.6     1
## 13 Perissodactyla 3.466667     2.9     4.4     3
## 14 Pilosa      14.400000    14.4    14.4     1
## 15 Primates    10.500000     8.0    17.0    12
## 16 Proboscidea  3.600000     3.3     3.9     2
## 17 Rodentia    12.468182     7.0    16.6    22
## 18 Scandentia   8.900000     8.9     8.9     1
## 19 Soricomorpha 11.100000     8.4    14.9     5
```